

Help on package xgboost:

NAME

xgboost - XGBoost: eXtreme Gradient Boosting library.

DESCRIPTION

Contributors: <https://github.com/dmlc/xgboost/blob/master/CONTRIBUTORS.md>

PACKAGE CONTENTS

compat
core
libpath
libxgboostwrapper
plotting
sklearn
training

CLASSES

```
builtins.object
  xgboost.core.Booster
  xgboost.core.DMatrix
sklearn.base.BaseEstimator(builtins.object)
  xgboost.sklearn.XGBModel
    xgboost.sklearn.XGBClassifier(xgboost.sklearn.XGBModel,
    sklearn.base.ClassifierMixin)
    xgboost.sklearn.XGBRegressor(xgboost.sklearn.XGBModel,
    sklearn.base.RegressorMixin)
```

```
class Booster(builtins.object)
|   "A Booster of of XGBoost.
|
|   Booster is the model of xgboost, that contains low level routines for
|   training, prediction and evaluation.
|
|   Methods defined here:
|
|   __copy__(self)
|
|   __deepcopy__(self)
|
|   __del__(self)
|
|   __getstate__(self)
|
|   __init__(self, params=None, cache=(), model_file=None)
|       Initialize the Booster.
|
|       Parameters
|       -----
|       params : dict
|           Parameters for boosters.
|       cache : list
|           List of cache items.
|       model_file : string
|           Path to the model file.
|
|   __setstate__(self, state)
```

```

boost(self, dtrain, grad, hess)
    Boost the booster for one iteration, with customized gradient statistics.

Parameters
-----
dtrain : DMatrix
    The training DMatrix.
grad : list
    The first order of gradient.
hess : list
    The second order of gradient.

copy(self)
    Copy the booster object.

Returns
-----
booster: `Booster`
    a copied booster model

dump_model(self, fout, fmap='', with_stats=False)
    Dump model into a text file.

Parameters
-----
foout : string
    Output file name.
fmap : string, optional
    Name of the file containing feature map names.
with_stats : bool (optional)
    Controls whether the split statistics are output.

eval(self, data, name='eval', iteration=0)
    Evaluate the model on mat.

Parameters
-----
data : DMatrix
    The dmatrix storing the input.

name : str, optional
    The name of the dataset.

iteration : int, optional
    The current iteration number.

Returns
-----
result: str
    Evaluation result string.

eval_set(self, evals, iteration=0, feval=None)
    Evaluate a set of data.

Parameters
-----

```

```

|     evals : list of tuples (DMatrix, string)
|           List of items to be evaluated.
|
|     iteration : int
|           Current iteration.
|
|     feval : function
|           Custom evaluation function.
|
|
| Returns
| -----
| result: str
|           Evaluation result string.
|
| get_dump(self, fmap='', with_stats=False)
|     Returns the dump the model as a list of strings.
|
| get_fscore(self, fmap='')
|     Get feature importance of each feature.
|
| Parameters
| -----
| fmap: str (optional)
|       The name of feature map file
|
| load_model(self, fname)
|     Load the model from a file.
|
| Parameters
| -----
| fname : string or a memory buffer
|       Input file name or memory buffer(see also save_raw)
|
| predict(self, data, output_margin=False, ntree_limit=0, pred_leaf=False)
|     Predict with data.
|
| NOTE: This function is not thread safe.
|       For each booster object, predict can only be called from one thread.
|       If you want to run prediction using multiple thread, call bst.copy() to
make copies
|       of model object and then call predict
|
| Parameters
| -----
| data : DMatrix
|       The dmatrix storing the input.
|
| output_margin : bool
|       Whether to output the raw untransformed margin value.
|
| ntree_limit : int
|       Limit number of trees in the prediction; defaults to 0 (use all trees).
|
| pred_leaf : bool
|       When this option is on, the output will be a matrix of (nsample, ntrees)
|       with each record indicating the predicted leaf index of each sample in each
tree.
|
|       Note that the leaf index of a tree is unique per tree, so you may find leaf 1
|       in both tree 1 and tree 0.

```

```

|     Returns
|     -----
|     prediction : numpy array
|
| save_model(self, fname)
|     Save the model to a file.
|
|     Parameters
|     -----
|     fname : string
|         Output file name
|
| save_raw(self)
|     Save the model to a in memory buffer represetation
|
|     Returns
|     -----
|     a in memory buffer represetation of the model
|
| set_param(self, params, value=None)
|     Set parameters into the Booster.
|
|     Parameters
|     -----
|     params: dict/list/str
|         list of key,value paris, dict of key to value or simply str key
|     value: optional
|         value of the specified parameter, when params is str key
|
| update(self, dtrain, iteration, fobj=None)
|     Update for one iteration, with objective function calculated internally.
|
|     Parameters
|     -----
|     dtrain : DMatrix
|         Training data.
|     iteration : int
|         Current iteration number.
|     fobj : function
|         Customized objective function.
|
| -----
| Data descriptors defined here:
|
| __dict__
|     dictionary for instance variables (if defined)
|
| __weakref__
|     list of weak references to the object (if defined)
|
| -----
| Data and other attributes defined here:
|
| feature_names = None
|
class DMatrix(builtins.object)

```

```

| Data Matrix used in XGBoost.
|
| DMatrix is a internal data structure that used by XGBoost
| which is optimized for both memory efficiency and training speed.
| You can construct DMatrix from numpy.arrays
|
| Methods defined here:
|
| __del__(self)
|
| __init__(self, data, label=None, missing=0.0, weight=None, silent=False,
feature_names=None, feature_types=None)
|     Data matrix used in XGBoost.
|
|     Parameters
|     -----
|     data : string/numpy array/scipy.sparse/pd.DataFrame
|           Data source of DMatrix.
|           When data is string type, it represents the path libsvm format txt file,
|           or binary file that xgboost can read from.
|     label : list or numpy 1-D array, optional
|           Label of the training data.
|     missing : float, optional
|           Value in the data which needs to be present as a missing value.
|     weight : list or numpy 1-D array , optional
|           Weight for each instance.
|     silent : boolean, optional
|           Whether print messages during construction
|     feature_names : list, optional
|           Set names for features.
|     feature_types : list, optional
|           Set types for features.
|
| get_base_margin(self)
|     Get the base margin of the DMatrix.
|
|     Returns
|     -----
|     base_margin : float
|
| get_float_info(self, field)
|     Get float property from the DMatrix.
|
|     Parameters
|     -----
|     field: str
|           The field name of the information
|
|     Returns
|     -----
|     info : array
|           a numpy array of float information of the data
|
| get_label(self)
|     Get the label of the DMatrix.
|
|     Returns

```

```

|         -----
|         label : array
|
| get_uint_info(self, field)
|     Get unsigned integer property from the DMatrix.
|
|     Parameters
|     -----
|     field: str
|         The field name of the information
|
|     Returns
|     -----
|     info : array
|         a numpy array of float information of the data
|
| get_weight(self)
|     Get the weight of the DMatrix.
|
|     Returns
|     -----
|     weight : array
|
| num_col(self)
|     Get the number of columns (features) in the DMatrix.
|
|     Returns
|     -----
|     number of columns : int
|
| num_row(self)
|     Get the number of rows in the DMatrix.
|
|     Returns
|     -----
|     number of rows : int
|
| save_binary(self, fname, silent=True)
|     Save DMatrix to an XGBoost buffer.
|
|     Parameters
|     -----
|     fname : string
|         Name of the output buffer file.
|     silent : bool (optional; default: True)
|         If set, the output is suppressed.
|
| set_base_margin(self, margin)
|     Set base margin of booster to start from.
|
|     This can be used to specify a prediction value of
|     existing model to be base_margin
|     However, remember margin is needed, instead of transformed prediction
|     e.g. for logistic regression: need to put in value before logistic transformation
|     see also example/demo.py
|
|     Parameters

```

```

|         -----
|         margin: array like
|             Prediction margin of each datapoint
|
| set_float_info(self, field, data)
|     Set float type property into the DMatrix.
|
|     Parameters
|     -----
|     field: str
|         The field name of the information
|
|     data: numpy array
|         The array of data to be set
|
| set_group(self, group)
|     Set group size of DMatrix (used for ranking).
|
|     Parameters
|     -----
|     group : array like
|         Group size of each group
|
| set_label(self, label)
|     Set label of dmatrix
|
|     Parameters
|     -----
|     label: array like
|         The label information to be set into DMatrix
|
| set_uint_info(self, field, data)
|     Set uint type property into the DMatrix.
|
|     Parameters
|     -----
|     field: str
|         The field name of the information
|
|     data: numpy array
|         The array of data to be set
|
| set_weight(self, weight)
|     Set weight of each instance.
|
|     Parameters
|     -----
|     weight : array like
|         Weight for each data point
|
| slice(self, rindex)
|     Slice the DMatrix and return a new DMatrix that only contains `rindex`.
|
|     Parameters
|     -----
|     rindex : list
|         List of indices to be selected.

```

```

|
|     Returns
|     -----
|     res : DMatrix
|         A new DMatrix containing only selected indices.
|
|     -----
| Data descriptors defined here:
|
|     __dict__
|         dictionary for instance variables (if defined)
|
|     __weakref__
|         list of weak references to the object (if defined)
|
|     feature_names
|         Get feature names (column labels).
|
|         Returns
|         -----
|         feature_names : list or None
|
|     feature_types
|         Get feature types (column types).
|
|         Returns
|         -----
|         feature_types : list or None
|
class XGBClassifier(XGBModel, sklearn.base.ClassifierMixin)
|     Implementation of the scikit-learn API for XGBoost classification.
|
|     Parameters
|     -----
|     max_depth : int
|         Maximum tree depth for base learners.
|     learning_rate : float
|         Boosting learning rate (xgb's "eta")
|     n_estimators : int
|         Number of boosted trees to fit.
|     silent : boolean
|         Whether to print messages while running boosting.
|     objective : string
|         Specify the learning task and the corresponding learning objective.
|
|     nthread : int
|         Number of parallel threads used to run xgboost.
|     gamma : float
|         Minimum loss reduction required to make a further partition on a leaf node of
the tree.
|     min_child_weight : int
|         Minimum sum of instance weight(hessian) needed in a child.
|     max_delta_step : int
|         Maximum delta step we allow each tree's weight estimation to be.
|     subsample : float
|         Subsample ratio of the training instance.
|     colsample_bytree : float

```



```

|     Subsample ratio of columns when constructing each tree.
| colsample_bylevel : float
|     Subsample ratio of columns for each split, in each level.
| reg_alpha : float (xgb's alpha)
|     L2 regularization term on weights
| reg_lambda : float (xgb's lambda)
|     L1 regularization term on weights
| scale_pos_weight : float
|     Balancing of positive and negative weights.
|
| base_score:
|     The initial prediction score of all instances, global bias.
| seed : int
|     Random number seed.
| missing : float, optional
|     Value in the data which needs to be present as a missing value. If
|     None, defaults to np.nan.
|
| Method resolution order:
|     XGBClassifier
|     XGBModel
|     sklearn.base.BaseEstimator
|     sklearn.base.ClassifierMixin
|     builtins.object
|
| Methods defined here:
|
| __init__(self, max_depth=3, learning_rate=0.1, n_estimators=100, silent=True,
objective='binary:logistic', nthread=-1, gamma=0, min_child_weight=1, max_delta_step=0,
subsample=1, colsample_bytree=1, colsample_bylevel=1, reg_alpha=0, reg_lambda=1,
scale_pos_weight=1, base_score=0.5, seed=0, missing=None)
|     Initialize self. See help(type(self)) for accurate signature.
|
| evals_result(self)
|     Return the evaluation results.
|
|     If eval_set is passed to the `fit` function, you can call evals_result() to
|     get evaluation results for all passed eval_sets. When eval_metric is also
|     passed to the `fit` function, the evals_result will contain the eval_metrics
|     passed to the `fit` function
|
| Returns
| -----
| evals_result : dictionary
|
| Example
| -----
| param_dist = {'objective':'binary:logistic', 'n_estimators':2}
|
| clf = xgb.XGBClassifier(**param_dist)
|
| clf.fit(X_train, y_train,
|         eval_set=[(X_train, y_train), (X_test, y_test)],
|         eval_metric='logloss',
|         verbose=True)
|
| evals_result = clf.evals_result()

```

```

|
|     The variable evals_result will contain:
|     {'validation_0': {'logloss': ['0.604835', '0.531479']},
|       'validation_1': {'logloss': ['0.41965', '0.17686']}}
|
|     fit(self, X, y, sample_weight=None, eval_set=None, eval_metric=None,
early_stopping_rounds=None, verbose=True)
|         Fit gradient boosting classifier
|
|         Parameters
|         -----
|         X : array_like
|             Feature matrix
|         y : array_like
|             Labels
|         sample_weight : array_like
|             Weight for each instance
|         eval_set : list, optional
|             A list of (X, y) pairs to use as a validation set for
|             early-stopping
|         eval_metric : str, callable, optional
|             If a str, should be a built-in evaluation metric to use. See
|             doc/parameter.md. If callable, a custom evaluation metric. The call
|             signature is func(y_predicted, y_true) where y_true will be a
|             DMatrix object such that you may need to call the get_label
|             method. It must return a str, value pair where the str is a name
|             for the evaluation and value is the value of the evaluation
|             function. This objective is always minimized.
|         early_stopping_rounds : int, optional
|             Activates early stopping. Validation error needs to decrease at
|             least every <early_stopping_rounds> round(s) to continue training.
|             Requires at least one item in evals. If there's more than one,
|             will use the last. Returns the model from the last iteration
|             (not the best one). If early stopping occurs, the model will
|             have two additional fields: bst.best_score and bst.best_iteration.
|         verbose : bool
|             If `verbose` and an evaluation set is used, writes the evaluation
|             metric measured on the validation set to stderr.
|
|     predict(self, data, output_margin=False, ntree_limit=0)
|
|     predict_proba(self, data, output_margin=False, ntree_limit=0)
|
|     -----
|     Methods inherited from XGBModel:
|
|     __setstate__(self, state)
|
|     booster(self)
|         Get the underlying xgboost Booster of this model.
|
|         This will raise an exception when fit was not called
|
|         Returns
|         -----
|         booster : a xgboost booster of underlying model

```

```

| get_params(self, deep=False)
|     Get parameter.s
|
| get_xgb_params(self)
|     Get xgboost type parameters.
|

```

```

| -----
| Methods inherited from sklearn.base.BaseEstimator:
|

```

```

| __repr__(self)
|     Return repr(self).
|

```

```

| set_params(self, **params)
|     Set the parameters of this estimator.
|

```

```

|     The method works on simple estimators as well as on nested objects
|     (such as pipelines). The former have parameters of the form
|     ``<component>__<parameter>`` so that it's possible to update each
|     component of a nested object.
|

```

```

| Returns
|

```

```

| -----
| self
|

```

```

| -----
| Data descriptors inherited from sklearn.base.BaseEstimator:
|

```

```

| __dict__
|     dictionary for instance variables (if defined)
|

```

```

| __weakref__
|     list of weak references to the object (if defined)
|

```

```

| -----
| Methods inherited from sklearn.base.ClassifierMixin:
|

```

```

| score(self, X, y, sample_weight=None)
|     Returns the mean accuracy on the given test data and labels.
|

```

```

|     In multi-label classification, this is the subset accuracy
|     which is a harsh metric since you require for each sample that
|     each label set be correctly predicted.
|

```

```

| Parameters
|

```

```

| -----

```

```

| X : array-like, shape = (n_samples, n_features)
|     Test samples.
|

```

```

| y : array-like, shape = (n_samples) or (n_samples, n_outputs)
|     True labels for X.
|

```

```

| sample_weight : array-like, shape = [n_samples], optional
|     Sample weights.
|

```

```

| Returns
|

```

```

| -----

```

```

| score : float
|

```

```

|           Mean accuracy of self.predict(X) wrt. y.

class XGBModel(sklearn.base.BaseEstimator)
|   Implementation of the Scikit-Learn API for XGBoost.
|
|   Parameters
|   -----
|   max_depth : int
|       Maximum tree depth for base learners.
|   learning_rate : float
|       Boosting learning rate (xgb's "eta")
|   n_estimators : int
|       Number of boosted trees to fit.
|   silent : boolean
|       Whether to print messages while running boosting.
|   objective : string
|       Specify the learning task and the corresponding learning objective.
|
|   nthread : int
|       Number of parallel threads used to run xgboost.
|   gamma : float
|       Minimum loss reduction required to make a further partition on a leaf node of
the tree.
|   min_child_weight : int
|       Minimum sum of instance weight(hessian) needed in a child.
|   max_delta_step : int
|       Maximum delta step we allow each tree's weight estimation to be.
|   subsample : float
|       Subsample ratio of the training instance.
|   colsample_bytree : float
|       Subsample ratio of columns when constructing each tree.
|   colsample_bylevel : float
|       Subsample ratio of columns for each split, in each level.
|   reg_alpha : float (xgb's alpha)
|       L2 regularization term on weights
|   reg_lambda : float (xgb's lambda)
|       L1 regularization term on weights
|   scale_pos_weight : float
|       Balancing of positive and negative weights.
|
|   base_score:
|       The initial prediction score of all instances, global bias.
|   seed : int
|       Random number seed.
|   missing : float, optional
|       Value in the data which needs to be present as a missing value. If
|       None, defaults to np.nan.
|
|   Method resolution order:
|       XGBModel
|       sklearn.base.BaseEstimator
|       builtins.object
|
|   Methods defined here:
|
|   __init__(self, max_depth=3, learning_rate=0.1, n_estimators=100, silent=True,
objective='reg:linear', nthread=-1, gamma=0, min_child_weight=1, max_delta_step=0,

```

```

subsample=1, colsample_bytree=1, colsample_bylevel=1, reg_alpha=0, reg_lambda=1,
scale_pos_weight=1, base_score=0.5, seed=0, missing=None)
|     Initialize self. See help(type(self)) for accurate signature.
|
| __setstate__(self, state)
|
| booster(self)
|     Get the underlying xgboost Booster of this model.
|
|     This will raise an exception when fit was not called
|
|     Returns
|     -----
|     booster : a xgboost booster of underlying model
|
| evals_result(self)
|     Return the evaluation results.
|
|     If eval_set is passed to the `fit` function, you can call evals_result() to
|     get evaluation results for all passed eval_sets. When eval_metric is also
|     passed to the `fit` function, the evals_result will contain the eval_metrics
|     passed to the `fit` function
|
|     Returns
|     -----
|     evals_result : dictionary
|
|     Example
|     -----
|     param_dist = {'objective':'binary:logistic', 'n_estimators':2}
|
|     clf = xgb.XGBModel(**param_dist)
|
|     clf.fit(X_train, y_train,
|             eval_set=[(X_train, y_train), (X_test, y_test)],
|             eval_metric='logloss',
|             verbose=True)
|
|     evals_result = clf.evals_result()
|
|     The variable evals_result will contain:
|     {'validation_0': {'logloss': ['0.604835', '0.531479']},
|      'validation_1': {'logloss': ['0.41965', '0.17686']}}
|
| fit(self, X, y, eval_set=None, eval_metric=None, early_stopping_rounds=None,
verbose=True)
|     Fit the gradient boosting model
|
|     Parameters
|     -----
|     X : array_like
|         Feature matrix
|     y : array_like
|         Labels
|     eval_set : list, optional
|         A list of (X, y) tuple pairs to use as a validation set for
|         early-stopping

```

```

|     eval_metric : str, callable, optional
|         If a str, should be a built-in evaluation metric to use. See
|         doc/parameter.md. If callable, a custom evaluation metric. The call
|         signature is func(y_predicted, y_true) where y_true will be a
|         DMatrix object such that you may need to call the get_label
|         method. It must return a str, value pair where the str is a name
|         for the evaluation and value is the value of the evaluation
|         function. This objective is always minimized.
|     early_stopping_rounds : int
|         Activates early stopping. Validation error needs to decrease at
|         least every <early_stopping_rounds> round(s) to continue training.
|         Requires at least one item in evals. If there's more than one,
|         will use the last. Returns the model from the last iteration
|         (not the best one). If early stopping occurs, the model will
|         have two additional fields: bst.best_score and bst.best_iteration.
|     verbose : bool
|         If `verbose` and an evaluation set is used, writes the evaluation
|         metric measured on the validation set to stderr.
|
|     get_params(self, deep=False)
|         Get parameter.s
|
|     get_xgb_params(self)
|         Get xgboost type parameters.
|
|     predict(self, data, output_margin=False, ntree_limit=0)
|
|     -----
|     Methods inherited from sklearn.base.BaseEstimator:
|
|     __repr__(self)
|         Return repr(self).
|
|     set_params(self, **params)
|         Set the parameters of this estimator.
|
|         The method works on simple estimators as well as on nested objects
|         (such as pipelines). The former have parameters of the form
|         ``<component>__<parameter>`` so that it's possible to update each
|         component of a nested object.
|
|         Returns
|         -----
|         self
|
|     -----
|     Data descriptors inherited from sklearn.base.BaseEstimator:
|
|     __dict__
|         dictionary for instance variables (if defined)
|
|     __weakref__
|         list of weak references to the object (if defined)
|
| class XGBRegressor(XGBModel, sklearn.base.RegressorMixin)
|     Implementation of the scikit-learn API for XGBoost regression.
|     Parameters

```

```

| -----
| max_depth : int
|     Maximum tree depth for base learners.
| learning_rate : float
|     Boosting learning rate (xgb's "eta")
| n_estimators : int
|     Number of boosted trees to fit.
| silent : boolean
|     Whether to print messages while running boosting.
| objective : string
|     Specify the learning task and the corresponding learning objective.
|
| nthread : int
|     Number of parallel threads used to run xgboost.
| gamma : float
|     Minimum loss reduction required to make a further partition on a leaf node of
the tree.
| min_child_weight : int
|     Minimum sum of instance weight(hessian) needed in a child.
| max_delta_step : int
|     Maximum delta step we allow each tree's weight estimation to be.
| subsample : float
|     Subsample ratio of the training instance.
| colsample_bytree : float
|     Subsample ratio of columns when constructing each tree.
| colsample_bylevel : float
|     Subsample ratio of columns for each split, in each level.
| reg_alpha : float (xgb's alpha)
|     L2 regularization term on weights
| reg_lambda : float (xgb's lambda)
|     L1 regularization term on weights
| scale_pos_weight : float
|     Balancing of positive and negative weights.
|
| base_score:
|     The initial prediction score of all instances, global bias.
| seed : int
|     Random number seed.
| missing : float, optional
|     Value in the data which needs to be present as a missing value. If
|     None, defaults to np.nan.
|
| Method resolution order:
|     XGBRegressor
|     XGBModel
|     sklearn.base.BaseEstimator
|     sklearn.base.RegressorMixin
|     builtins.object
|
| Methods inherited from XGBModel:
|
| __init__(self, max_depth=3, learning_rate=0.1, n_estimators=100, silent=True,
objective='reg:linear', nthread=-1, gamma=0, min_child_weight=1, max_delta_step=0,
subsample=1, colsample_bytree=1, colsample_bylevel=1, reg_alpha=0, reg_lambda=1,
scale_pos_weight=1, base_score=0.5, seed=0, missing=None)
|     Initialize self. See help(type(self)) for accurate signature.
|

```

```

|   __setstate__(self, state)
|
|   booster(self)
|       Get the underlying xgboost Booster of this model.
|
|       This will raise an exception when fit was not called
|
|       Returns
|       -----
|       booster : a xgboost booster of underlying model
|
|   evals_result(self)
|       Return the evaluation results.
|
|       If eval_set is passed to the `fit` function, you can call evals_result() to
|       get evaluation results for all passed eval_sets. When eval_metric is also
|       passed to the `fit` function, the evals_result will contain the eval_metrics
|       passed to the `fit` function
|
|       Returns
|       -----
|       evals_result : dictionary
|
|       Example
|       -----
|       param_dist = {'objective':'binary:logistic', 'n_estimators':2}
|
|       clf = xgb.XGBModel(**param_dist)
|
|       clf.fit(X_train, y_train,
|               eval_set=[(X_train, y_train), (X_test, y_test)],
|               eval_metric='logloss',
|               verbose=True)
|
|       evals_result = clf.evals_result()
|
|       The variable evals_result will contain:
|       {'validation_0': {'logloss': ['0.604835', '0.531479']},
|        'validation_1': {'logloss': ['0.41965', '0.17686']}}
|
|   fit(self, X, y, eval_set=None, eval_metric=None, early_stopping_rounds=None,
|       verbose=True)
|       Fit the gradient boosting model
|
|       Parameters
|       -----
|       X : array_like
|           Feature matrix
|       y : array_like
|           Labels
|       eval_set : list, optional
|           A list of (X, y) tuple pairs to use as a validation set for
|           early-stopping
|       eval_metric : str, callable, optional
|           If a str, should be a built-in evaluation metric to use. See
|           doc/parameter.md. If callable, a custom evaluation metric. The call
|           signature is func(y_predicted, y_true) where y_true will be a

```



```

|         DMatrix object such that you may need to call the get_label
|         method. It must return a str, value pair where the str is a name
|         for the evaluation and value is the value of the evaluation
|         function. This objective is always minimized.
|     early_stopping_rounds : int
|         Activates early stopping. Validation error needs to decrease at
|         least every <early_stopping_rounds> round(s) to continue training.
|         Requires at least one item in evals. If there's more than one,
|         will use the last. Returns the model from the last iteration
|         (not the best one). If early stopping occurs, the model will
|         have two additional fields: bst.best_score and bst.best_iteration.
|     verbose : bool
|         If `verbose` and an evaluation set is used, writes the evaluation
|         metric measured on the validation set to stderr.
|
|     get_params(self, deep=False)
|         Get parameter.s
|
|     get_xgb_params(self)
|         Get xgboost type parameters.
|
|     predict(self, data, output_margin=False, ntree_limit=0)
|
|     -----
|     Methods inherited from sklearn.base.BaseEstimator:
|
|     __repr__(self)
|         Return repr(self).
|
|     set_params(self, **params)
|         Set the parameters of this estimator.
|
|         The method works on simple estimators as well as on nested objects
|         (such as pipelines). The former have parameters of the form
|         ``<component>__<parameter>`` so that it's possible to update each
|         component of a nested object.
|
|         Returns
|         -----
|         self
|
|     -----
|     Data descriptors inherited from sklearn.base.BaseEstimator:
|
|     __dict__
|         dictionary for instance variables (if defined)
|
|     __weakref__
|         list of weak references to the object (if defined)
|
|     -----
|     Methods inherited from sklearn.base.RegressorMixin:
|
|     score(self, X, y, sample_weight=None)
|         Returns the coefficient of determination R^2 of the prediction.
|
|         The coefficient R^2 is defined as (1 - u/v), where u is the regression

```

```

|         sum of squares ((y_true - y_pred) ** 2).sum() and v is the residual
|         sum of squares ((y_true - y_true.mean()) ** 2).sum().
|         Best possible score is 1.0 and it can be negative (because the
|         model can be arbitrarily worse). A constant model that always
|         predicts the expected value of y, disregarding the input features,
|         would get a R^2 score of 0.0.

```

Parameters

```

|         X : array-like, shape = (n_samples, n_features)

```

```

|             Test samples.

```

```

|         y : array-like, shape = (n_samples) or (n_samples, n_outputs)
|             True values for X.

```

```

|         sample_weight : array-like, shape = [n_samples], optional
|             Sample weights.

```

Returns

```

|         score : float
|             R^2 of self.predict(X) wrt. y.

```

FUNCTIONS

```

cv(params, dtrain, num_boost_round=10, nfold=3, metrics=(), obj=None, feval=None,
maximize=False, early_stopping_rounds=None, fpreproc=None, as_pandas=True,
show_progress=None, show_stdv=True, seed=0)
    Cross-validation with given paramaters.

```

Parameters

```

|         params : dict

```

```

|             Booster params.

```

```

|         dtrain : DMatrix

```

```

|             Data to be trained.

```

```

|         num_boost_round : int

```

```

|             Number of boosting iterations.

```

```

|         nfold : int

```

```

|             Number of folds in CV.

```

```

|         metrics : list of strings

```

```

|             Evaluation metrics to be watched in CV.

```

```

|         obj : function

```

```

|             Custom objective function.

```

```

|         feval : function

```

```

|             Custom evaluation function.

```

```

|         maximize : bool

```

```

|             Whether to maximize feval.

```

```

|         early_stopping_rounds: int

```

```

|             Activates early stopping. CV error needs to decrease at least
|             every <early_stopping_rounds> round(s) to continue.

```

```

|             Last entry in evaluation history is the one from best iteration.

```

```

|         fpreproc : function

```

```

|             Preprocessing function that takes (dtrain, dtest, param) and returns
|             transformed versions of those.

```

```

|         as_pandas : bool, default True

```

```

|             Return pd.DataFrame when pandas is installed.

```

```

|             If False or pandas is not installed, return np.ndarray

```

```

show_progress : bool or None, default None
    Whether to display the progress. If None, progress will be displayed
    when np.ndarray is returned.
show_stdv : bool, default True
    Whether to display the standard deviation in progress.
    Results are not affected, and always contains std.
seed : int
    Seed used to generate the folds (passed to numpy.random.seed).

```

Returns

```
evaluation history : list(string)
```

```

plot_importance(booster, ax=None, height=0.2, xlim=None, ylim=None, title='Feature
importance', xlabel='F score', ylabel='Features', grid=True, **kwargs)
    Plot importance based on fitted trees.

```

Parameters

```

booster : Booster, XGBModel or dict
    Booster or XGBModel instance, or dict taken by Booster.get_fscore()
ax : matplotlib Axes, default None
    Target axes instance. If None, new figure and axes will be created.
height : float, default 0.2
    Bar height, passed to ax.barh()
xlim : tuple, default None
    Tuple passed to axes.xlim()
ylim : tuple, default None
    Tuple passed to axes.ylim()
title : str, default "Feature importance"
    Axes title. To disable, pass None.
xlabel : str, default "F score"
    X axis title label. To disable, pass None.
ylabel : str, default "Features"
    Y axis title label. To disable, pass None.
kwargs :
    Other keywords passed to ax.barh()

```

Returns

```
ax : matplotlib Axes
```

```

plot_tree(booster, num_trees=0, rankdir='UT', ax=None, **kwargs)
    Plot specified tree.

```

Parameters

```

booster : Booster, XGBModel
    Booster or XGBModel instance
num_trees : int, default 0
    Specify the ordinal number of target tree
rankdir : str, default "UT"
    Passed to graphviz via graph_attr
ax : matplotlib Axes, default None
    Target axes instance. If None, new figure and axes will be created.
kwargs :
    Other keywords passed to to_graphviz

```

Returns

ax : matplotlib Axes

to_graphviz(booster, num_trees=0, rankdir='UT', yes_color='#0000FF', no_color='#FF0000', **kwargs)

Convert specified tree to graphviz instance. IPython can automatically plot the returned graphviz instance. Otherwise, you should call .render() method of the returned graphviz instance.

Parameters

booster : Booster, XGBModel

Booster or XGBModel instance

num_trees : int, default 0

Specify the ordinal number of target tree

rankdir : str, default "UT"

Passed to graphviz via graph_attr

yes_color : str, default '#0000FF'

Edge color when meets the node condition.

no_color : str, default '#FF0000'

Edge color when doesn't meet the node condition.

kwargs :

Other keywords passed to graphviz graph_attr

Returns

ax : matplotlib Axes

train(params, dtrain, num_boost_round=10, evals=(), obj=None, feval=None, maximize=False, early_stopping_rounds=None, evals_result=None, verbose_eval=True, learning_rates=None, xgb_model=None)

Train a booster with given parameters.

Parameters

params : dict

Booster params.

dtrain : DMatrix

Data to be trained.

num_boost_round: int

Number of boosting iterations.

watchlist (evals): list of pairs (DMatrix, string)

List of items to be evaluated during training, this allows user to watch performance on the validation set.

obj : function

Customized objective function.

feval : function

Customized evaluation function.

maximize : bool

Whether to maximize feval.

early_stopping_rounds: int

Activates early stopping. Validation error needs to decrease at least every <early_stopping_rounds> round(s) to continue training.

Requires at least one item in evals.

If there's more than one, will use the last.

Returns the model from the last iteration (not the best one).
 If early stopping occurs, the model will have three additional fields:
 bst.best_score, bst.best_iteration and bst.best_ntree_limit.
 (Use bst.best_ntree_limit to get the correct value if num_parallel_tree
 and/or num_class appears in the parameters)

evals_result: dict

This dictionary stores the evaluation results of all the items in watchlist.
 Example: with a watchlist containing [(dtest,'eval'), (dtrain,'train')] and
 and a paramater containing ('eval_metric', 'logloss')

Returns: {'train': {'logloss': ['0.48253', '0.35953']},
 'eval': {'logloss': ['0.480385', '0.357756']}}

verbose_eval : bool or int

Requires at least one item in evals.

If `verbose_eval` is True then the evaluation metric on the validation set is
 printed at each boosting stage.

If `verbose_eval` is an integer then the evaluation metric on the validation set
 is printed at every given `verbose_eval` boosting stage. The last boosting stage
 / the boosting stage found by using `early_stopping_rounds` is also printed.

Example: with verbose_eval=4 and at least one item in evals, an evaluation metric
 is printed every 4 boosting stages, instead of every boosting stage.

learning_rates: list or function

List of learning rate for each boosting round
 or a customized function that calculates eta in terms of
 current number of round and the total number of boosting round (e.g. yields
 learning rate decay)

- list l: eta = l[boosting round]

- function f: eta = f(boosting round, num_boost_round)

xgb_model : file name of stored xgb model or 'Booster' instance

Xgb model to be loaded before training (allows training continuation).

Returns

booster : a trained booster model

DATA

```
__all__ = ['DMatrix', 'Booster', 'train', 'cv', 'XGBModel', 'XGBClassi...
__warningregistry__ = {'version': 207, ("unclosed file <_io.TextIOWrap...
```

VERSION

0.4

FILE

```
/home/shuangyangwang/anaconda3/lib/python3.5/site-packages/xgboost-0.4-py3.5.egg/xgboost/_
_init__.py
```